Isha Patel
0801IT221065

Assignment - 6

Q1) Run 40 commands of Hive on Hive terminal.

Start Hive and connect -

1. Run 'hive' and when prompted with beeline terminal, enter '!connect jdbc:hive2://'
2. At the end we see the hive terminal connected successfully.

```
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar: file:/home/manasvi_ajmera/Documents/SGSITS/SEM
-slf4j-impl-2.18.0.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar: file:/home/manasvi_ajmera/Documents/SGSITS/SEM hare/hadoop/common/lib/slf4j-
reload4j-1.7.36.jar!/org/slf4j/impl/StaticLogg
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explana
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFacto
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar: file:/home/manasvi_ajmera/Documents/SGSITS/SEM
-slf4j-impl-2.18.0.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar: file:/home/manasvi_ajmera/Documents/SGSITS/SEM
hare/hadoop/common/lib/slf4j-reload4j-1.7.36.jar!/org/slf4j/impl/StaticLogg
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explana SLF4J: Actual binding is of type
[org.apache.logging.slf4j.Log4jLoggerFacto
Beeline version 4.0.0 by Apache Hive
beeline> !connect jdbc:hive2://
Connecting to jdbc:hive2://
Enter username for jdbc:hive2://:
Enter password for jdbc:hive2://:
```

```
Connected to: Apache Hive (version 4.0.0)
Driver: Hive JDBC (version 4.0.0)
Transaction isolation: TRANSACTION_REPEATABLE_READ
0: jdbc:hive2://>
```

Hive Commands -

1. SHOW DATABASES;

```
0: jdbc:hive2://> SHOW DATABASES;
+----------------+
| database_name  |
+----------------+
| default        |
+----------------+
1 row selected (0.703 seconds)
```

2. CREATE DATABASE IF NOT EXISTS bda_assign;

Isha Patel
0801IT221065

```
0: jdbc:hive2://> CREATE DATABASE IF NOT EXISTS bda_assign;
No rows affected (0.047 seconds)
```

3. DESCRIBE DATABASE bda_assign;

```
0: jdbc:hive2://> DESCRIBE DATABASE bda_assign;
26/03/25 03:06:28 [451acfe7-4ff5-4045-ba0e-9fe5bce20c80 main]: WARN lazy.LazyStruct: Extra bytes detected at the end of the row! Ignoring similar problems.
+------------+---------+-----------------------------------+-------------------------------------+----------------+------------+----------------+---------------+
|   db_name  | comment |              location             |            managedlocation          |   owner_name   | owner_type | connector_name | remote_dbname |
+------------+---------+-----------------------------------+-------------------------------------+----------------+------------+----------------+---------------+

| bda_assign |         | hdfs://localhost:9000/user/hive/warehouse/external/bda_assign/bda_assign.db | manasvi-ajmera | USER       |                |               |
+------------+---------+-----------------------------------+-------------------------------------+----------------+------------+----------------+---------------+
1 row selected (0.025 seconds)
```

4. USE bda_assign;

```
0: jdbc:hive2://> USE bda_assign;
No rows affected (0.019 seconds)
```

5. SELECT current_database();

```
0: jdbc:hive2://> SELECT current_database();
WARNING: An illegal reflective access operation has occurred
WARNING: Illegal reflective access by org.apache.calcite.avatica.com.google.protobuf. UnsafeUtil (file:/home/manasvi-ajmera/Documents/SGSITS/SEM_
8/BDA/Assignment_6/hive/lib/avatica-1.12.0.jar) to field java.nio.Buffer.a ddress
WARNING: Please consider reporting this to the maintainers of org.apache.calcite.avatica.com.google.protobuf. UnsafeUtil
WARNING: Use--illegal-access=warn to enable warnings of further illegal reflective access operations
WARNING: All illegal access operations will be denied in a future release
26/03/25 03:09:19 [451acfe7-4ff5-4045-ba0e-9fe5bce20c80 main]: WARN optimizer. SimpleFetchOptimizer: Cannot de
termine basic stats for table: _dummy_database@_dummy_table from metastore. Falling back.
```

6. CREATE TABLE employees (emp_id INT, name STRING, department STRING, salary DOUBLE);

```
0: jdbc:hive2://> CREATE TABLE employees (emp_id INT, name STRING, department STRING, salary DOUBLE);
No rows affected (0.134 seconds)
```

7. SHOW TABLES;

```
0: jdbc:hive2://> SHOW TABLES;
+-------------+
|  tab_name   |
+-------------+
| employees   |
+-------------+
1 row selected (0.068 seconds)
```

8. DESCRIBE employees;

```
0: jdbc:hive2://> DESCRIBE employees;
+--------------+------------+----------+
|  col_name    | data_type  | comment  |
+--------------+------------+----------+
| emp_id       | int        |          |
| name         | string     |          |
| department   | string     |          |
| salary       | double     |          |
+--------------+------------+----------+
4 rows selected (0.07 seconds)
```

9. DESCRIBE FORMATTED employees;

Isha Patel
0801IT221065

```
0: jdbc:hive2://> DESCRIBE FORMATTED employees;
+-------------------------------+----------------------------------------------------+-----------------------------------------------------
|           col_name            |                     data_type                      |                       comment
+-------------------------------+----------------------------------------------------+-----------------------------------------------------
| emp_id                        | int                                                |
| name                          | string                                             |
| department                    | string                                             |
| salary                        | double                                             |
|                               | NULL                                               | NULL
| # Detailed Table Information  | NULL                                               | NULL
| Database:                     | bda_assign                                         | NULL
| OwnerType:                    | USER                                               | NULL
| Owner:                        | abhishek-verma                                     | NULL
| CreateTime:                   | Wed Mar 25 03:12:16 UTC 2026                       | NULL
| LastAccessTime:               | UNKNOWN                                            | NULL
| Retention:                    | 0                                                  | NULL
| Location:                     | hdfs://localhost:9000/user/hive/warehouse/external/bda_assign.db/employees | NULL
                      |
| Table Type:                   | EXTERNAL_TABLE                                     | NULL
| Table Parameters:             | NULL                                               | NULL
|                               | COLUMN_STATS_ACCURATE                              | {\"BASIC_STATS\":\"true\",\"COLUMN_STATS\":{\"depa
ment\":\"true\",\"emp_id\":\"true\",\"name\":\"true\",\"salary\":\"true\"}} |
|                               | EXTERNAL                                           | TRUE
|                               | TRANSLATED_TO_EXTERNAL                             | TRUE
|                               | bucketing_version                                  | 2
|                               | external.table.purge                               | TRUE
|                               | numFiles                                           | 0
|                               | numRows                                            | 0
|                               | rawDataSize                                        | 0
|                               | totalSize                                          | 0
|                               | transient_lastDdlTime                              | 1774408336
|                               | NULL                                               | NULL
| # Storage Information         | NULL                                               | NULL
| SerDe Library:                | org.apache.hadoop.hive.serde2.lazy.LazySimpleSerDe | NULL
| InputFormat:                  | org.apache.hadoop.mapred.TextInputFormat           | NULL
| OutputFormat:                 | org.apache.hadoop.hive.ql.io.HiveIgnoreKeyTextOutputFormat | NULL
      |
| Compressed:                   | No                                                 | NULL
| Num Buckets:                  | -1                                                 | NULL
| Bucket Columns:               | []                                                 | NULL
| Sort Columns:                 | []                                                 | NULL
| Storage Desc Params:          | NULL                                               | NULL
|                               | serialization.format                               | 1
+-------------------------------+----------------------------------------------------+-----------------------------------------------------
36 rows selected (0.056 seconds)
```

10. ALTER TABLE employees RENAME TO staff_members;

```
0: jdbc:hive2://> ALTER TABLE employees RENAME TO staff_members;
26/03/25 03:17:31 [HiveServer2-Background-Pool: Thread-163]: WAR
N metadata.Hive: Cannot get a table snapshot for employees
No rows affected (0.098 seconds)
0: jdbc:hive2://> SHOW TABLES;
+----------------+
|    tab_name    |
+----------------+
| staff_members  |
+----------------+
1 row selected (0.028 seconds)
```

11. ALTER TABLE staff_members ADD COLUMNS (join_date STRING);

```
0: jdbc:hive2://> ALTER TABLE staff_members ADD COLUMNS (join_date STRING);
26/03/25 03:18:57 [HiveServer2-Background-Pool: Thread-174]: WARN metadata.Hi
ve: Cannot get a table snapshot for staff_members
No rows affected (0.075 seconds)
0: jdbc:hive2://> DESCRIBE staff_members;
+-------------+------------+----------+
|  col_name   | data_type  | comment  |
+-------------+------------+----------+
| emp_id      | int        |          |
| name        | string     |          |
| department  | string     |          |
| salary      | double     |          |
| join_date   | string     |          |
+-------------+------------+----------+
5 rows selected (0.03 seconds)
```

12. ALTER TABLE staff_members CHANGE COLUMN name full_name STRING;

```
0: jdbc:hive2://> ALTER TABLE staff_members CHANGE COLUMN name full_name STRING;
26/03/25 03:20:19 [HiveServer2-Background-Pool: Thread-182]: WARN metadata.Hive:
 Cannot get a table snapshot for staff_members
No rows affected (0.043 seconds)
0: jdbc:hive2://> DESCRIBE staff_members;
+-------------+------------+----------+
|  col_name   | data_type  | comment  |
+-------------+------------+----------+
| emp_id      | int        |          |
| full_name   | string     |          |
| department  | string     |          |
| salary      | double     |          |
| join_date   | string     |          |
+-------------+------------+----------+
5 rows selected (0.026 seconds)
```

13. CREATE TABLE contractors (id INT, name STRING);

```
0: jdbc:hive2://> CREATE TABLE contractors (id INT, name STRING);
No rows affected (0.043 seconds)
0: jdbc:hive2://> SHOW TABLES;
+----------------+
|    tab_name    |
+----------------+
| contractors    |
| staff_members  |
+----------------+
2 rows selected (0.029 seconds)
```

14. SHOW COLUMNS IN staff_members;

```
0: jdbc:hive2://> SHOW COLUMNS IN staff_members;
+-------------+
|    field    |
+-------------+
| emp_id      |
| full_name   |
| department  |
| salary      |
| join_date   |
+-------------+
5 rows selected (0.032 seconds)
```

15. TRUNCATE TABLE contractors;

```
0: jdbc:hive2://> TRUNCATE TABLE contractors;
No rows affected (0.101 seconds)
```

16. INSERT INTO staff_members VALUES (101, 'Manasvi', 'Engineering', 85000.0, '2023-01-15');

```
MapReduce Jobs Launched:
Stage-Stage-1:  Cumulative CPU: 0.86 sec   HDFS Read: 0 HDFS Write: 169356285 HDFS EC Read: 0 SUCCESS
Total MapReduce CPU Time Spent: 860 msec
1 row affected (1.572 seconds)
```

17. INSERT INTO staff_members VALUES (102, 'Surendra', 'Marketing', 65000.0, '2023-03-22'), (103, 'Akshat', 'Engineering', 92000.0, '2022-11-05'), (104, 'Anjali', 'HR', 75000.0, '2024-02-10');

```
MapReduce Jobs Launched:
Stage-Stage-1:  Cumulative CPU: 0.49 sec   HDFS Read: 806 HDFS Write: 169357369 HDFS EC Read: 0 SUCCES
Total MapReduce CPU Time Spent: 490 msec
3 rows affected (1.401 seconds)
```

18. INSERT INTO contractors VALUES (1, 'Manasvi');

```
MapReduce Jobs Launched:
Stage-Stage-1:  Cumulative CPU: 0.7 sec   HDFS Read: 1782 HDFS Write: 169358114 HDFS EC Read: 0 SUCCES
Total MapReduce CPU Time Spent: 700 msec
1 row affected (1.769 seconds)
```

19. CREATE TABLE engineering_staff AS SELECT * FROM staff_members WHERE department = 'Engineering';

```
MapReduce Jobs Launched:
Stage-Stage-1:  Cumulative CPU: 0.66 sec   HDFS Read: 2580 HDFS Write: 169358935 HDFS EC Read: 0 SUCCE
Total MapReduce CPU Time Spent: 660 msec
2 rows affected (2.842 seconds)
```

20. SELECT * FROM staff_members;

Isha Patel
0801IT221065

```
0: jdbc:hive2://> SELECT * FROM staff_members;
26/03/25 04:00:26 [451acfe7-4ff5-4045-ba0e-9fe5bce20c80 main]: WARN calcite.RelOptHiveTable: No Stats for bda_assign@staff_members, C
olumns: full_name, join_date
No Stats for bda_assign@staff_members, Columns: full_name, join_date
26/03/25 04:00:26 [451acfe7-4ff5-4045-ba0e-9fe5bce20c80 main]: WARN optimizer.SimpleFetchOptimizer: Table bda_assign@staff_members is
 external table, falling back to filesystem scan.
+-----------------------+-------------------------+---------------------------+-----------------------+-------------------------+
| staff_members.emp_id  | staff_members.full_name | staff_members.department  | staff_members.salary  | staff_members.join_date |
+-----------------------+-------------------------+---------------------------+-----------------------+-------------------------+
| 101                   | Abhishek                | Engineering               | 85000.0               | 2023-01-15              |
| 102                   | Surendra                | Marketing                 | 65000.0               | 2023-03-22              |
| 103                   | Akshat                  | Engineering               | 92000.0               | 2022-11-05              |
| 104                   | Anjali                  | HR                        | 75000.0               | 2024-02-10              |
+-----------------------+-------------------------+---------------------------+-----------------------+-------------------------+
4 rows selected (0.054 seconds)
```

21. SELECT full_name, salary FROM staff_members;

```
0: jdbc:hive2://> SELECT full_name, salary FROM staff_members;
26/03/25 04:00:55 [451acfe7-4ff5-4045-ba0e-9fe5bce20c80 main]:
olumns: full_name
No Stats for bda_assign@staff_members, Columns: full_name
26/03/25 04:00:55 [451acfe7-4ff5-4045-ba0e-9fe5bce20c80 main]:
 external table, falling back to filesystem scan.
+-------------+----------+
| full_name   |  salary  |
+-------------+----------+
| Abhishek    | 85000.0  |
| Surendra    | 65000.0  |
| Akshat      | 92000.0  |
| Anjali      | 75000.0  |
+-------------+----------+
4 rows selected (0.05 seconds)
```

22. SELECT * FROM staff_members WHERE salary > 70000;

```
0: jdbc:hive2://> SELECT * FROM staff_members WHERE salary > 70000;
26/03/25 04:01:22 [451acfe7-4ff5-4045-ba0e-9fe5bce20c80 main]: WARN calcite.RelOptHiveTable: No Stats for bda_assign@staff_members, C
olumns: full_name, join_date
No Stats for bda_assign@staff_members, Columns: full_name, join_date
26/03/25 04:01:22 [451acfe7-4ff5-4045-ba0e-9fe5bce20c80 main]: WARN optimizer.SimpleFetchOptimizer: Table bda_assign@staff_members is
 external table, falling back to filesystem scan.
+-----------------------+-------------------------+---------------------------+-----------------------+-------------------------+
| staff_members.emp_id  | staff_members.full_name | staff_members.department  | staff_members.salary  | staff_members.join_date |
+-----------------------+-------------------------+---------------------------+-----------------------+-------------------------+
| 101                   | Abhishek                | Engineering               | 85000.0               | 2023-01-15              |
| 103                   | Akshat                  | Engineering               | 92000.0               | 2022-11-05              |
| 104                   | Anjali                  | HR                        | 75000.0               | 2024-02-10              |
+-----------------------+-------------------------+---------------------------+-----------------------+-------------------------+
3 rows selected (0.073 seconds)
```

23. SELECT * FROM staff_members ORDER BY salary DESC;

Isha Patel
0801IT221065

```
MapReduce Total cumulative CPU time: 470 msec
Ended Job = job_local959925271_0005
MapReduce Jobs Launched:
Stage-Stage-1:  Cumulative CPU: 0.47 sec   HDFS Read: 4674 HDFS Write: 169359550 HDFS EC Read: 0 SUCCESS
Total MapReduce CPU Time Spent: 470 msec
+-----------------------+--------------------------+---------------------------+-----------------------+--------------------------+
| staff_members.emp_id  | staff_members.full_name  | staff_members.department  | staff_members.salary  | staff_members.join_date  |
+-----------------------+--------------------------+---------------------------+-----------------------+--------------------------+
| 103                   | Akshat                   | Engineering               | 92000.0               | 2022-11-05               |
| 101                   | Abhishek                 | Engineering               | 85000.0               | 2023-01-15               |
| 104                   | Anjali                   | HR                        | 75000.0               | 2024-02-10               |
| 102                   | Surendra                 | Marketing                 | 65000.0               | 2023-03-22               |
+-----------------------+--------------------------+---------------------------+-----------------------+--------------------------+
4 rows selected (1.256 seconds)
```

24. SELECT * FROM staff_members LIMIT 2;

```
0: jdbc:hive2://> SELECT * FROM staff_members LIMIT 2;
26/03/25 04:02:47 [451acfe7-4ff5-4045-ba0e-9fe5bce20c80 main]: WARN calcite.RelOptHiveTable: No Stats for bda_assign@staff_members,
olumns: full_name, join_date
No Stats for bda_assign@staff_members, Columns: full_name, join_date
26/03/25 04:02:47 [451acfe7-4ff5-4045-ba0e-9fe5bce20c80 main]: WARN optimizer.SimpleFetchOptimizer: Table bda_assign@staff_members i
 external table, falling back to filesystem scan.
+-----------------------+--------------------------+---------------------------+-----------------------+--------------------------+
| staff_members.emp_id  | staff_members.full_name  | staff_members.department  | staff_members.salary  | staff_members.join_date  |
+-----------------------+--------------------------+---------------------------+-----------------------+--------------------------+
| 101                   | Abhishek                 | Engineering               | 85000.0               | 2023-01-15               |
| 102                   | Surendra                 | Marketing                 | 65000.0               | 2023-03-22               |
+-----------------------+--------------------------+---------------------------+-----------------------+--------------------------+
2 rows selected (0.057 seconds)
```

25. SELECT DISTINCT department FROM staff_members;

```
MapReduce Jobs Launched:
Stage-Stage-1:  Cumulative CPU: 0.38 sec
HDFS Read: 5910 HDFS Write: 169359994 HDFS
EC Read: 0 SUCCESS
Total MapReduce CPU Time Spent: 380 msec
+--------------+
|  department  |
+--------------+
| Engineering  |
| HR           |
| Marketing    |
+--------------+
3 rows selected (1.289 seconds)
```

26. SELECT * FROM staff_members WHERE department = 'HR' OR salary > 90000;

```
0: jdbc:hive2://> SELECT * FROM staff_members WHERE department = 'HR' OR salary > 90000;
26/03/25 04:06:27 [451acfe7-4ff5-4045-ba0e-9fe5bce20c80 main]: WARN calcite.RelOptHiveTable: No Stats for bda_assign@staff_members,
umns: full_name, join_date
No Stats for bda_assign@staff_members, Columns: full_name, join_date
26/03/25 04:06:27 [451acfe7-4ff5-4045-ba0e-9fe5bce20c80 main]: WARN optimizer.SimpleFetchOptimizer: Table bda_assign@staff_members i
xternal table, falling back to filesystem scan.
+-----------------------+--------------------------+---------------------------+-----------------------+--------------------------+
| staff_members.emp_id  | staff_members.full_name  | staff_members.department  | staff_members.salary  | staff_members.join_date  |
+-----------------------+--------------------------+---------------------------+-----------------------+--------------------------+
| 103                   | Akshat                   | Engineering               | 92000.0               | 2022-11-05               |
| 104                   | Anjali                   | HR                        | 75000.0               | 2024-02-10               |
+-----------------------+--------------------------+---------------------------+-----------------------+--------------------------+
2 rows selected (0.093 seconds)
```

Isha Patel
0801IT221065

27. SELECT * FROM staff_members WHERE full_name LIKE 'A%';

```
0: jdbc:hive2://> SELECT * FROM staff_members WHERE full_name LIKE 'A%';
26/03/25 04:06:55 [451acfe7-4ff5-4045-ba0e-9fe5bce20c80 main]: WARN calcite.RelOptHiveTable: No Stats for bda_assign@staff_members,
umns: full_name, join_date
No Stats for bda_assign@staff_members, Columns: full_name, join_date
26/03/25 04:06:55 [451acfe7-4ff5-4045-ba0e-9fe5bce20c80 main]: WARN optimizer.SimpleFetchOptimizer: Table bda_assign@staff_members i
xternal table, falling back to filesystem scan.
+----------------------+-------------------------+--------------------------+----------------------+-------------------------+
| staff_members.emp_id | staff_members.full_name | staff_members.department | staff_members.salary | staff_members.join_date |
+----------------------+-------------------------+--------------------------+----------------------+-------------------------+
| 101                  | Abhishek                | Engineering              | 85000.0              | 2023-01-15              |
| 103                  | Akshat                  | Engineering              | 92000.0              | 2022-11-05              |
| 104                  | Anjali                  | HR                       | 75000.0              | 2024-02-10              |
+----------------------+-------------------------+--------------------------+----------------------+-------------------------+
3 rows selected (0.089 seconds)
```

28. SELECT COUNT(*) FROM staff_members;

```
MapReduce Jobs Launched:
Stage-Stage-1:  Cumulative CPU: 0.64 sec
HDFS Read: 7172 HDFS Write: 169360243 HDFS
EC Read: 0 SUCCESS
Total MapReduce CPU Time Spent: 640 msec
+------+
| _c0  |
+------+
| 4    |
+------+
1 row selected (1.237 seconds)
```

29. SELECT department, COUNT(*) FROM staff_members GROUP BY department;

```
MapReduce Jobs Launched:
Stage-Stage-1:  Cumulative CPU: 0.31 sec
HDFS Read: 7696 HDFS Write: 169360498 HDFS
EC Read: 0 SUCCESS
Total MapReduce CPU Time Spent: 310 msec
+--------------+------+
|  department  | _c1  |
+--------------+------+
| Engineering  | 2    |
| HR           | 1    |
| Marketing    | 1    |
+--------------+------+
3 rows selected (1.241 seconds)
```

30. SELECT department, AVG(salary) FROM staff_members GROUP BY department;

```
MapReduce Jobs Launched:
Stage-Stage-1:  Cumulative CPU: 0.19 sec
HDFS Read: 8326 HDFS Write: 169360824 HDF
EC Read: 0 SUCCESS
Total MapReduce CPU Time Spent: 190 msec
+--------------+----------+
|  department  |   _c1    |
+--------------+----------+
| Engineering  | 88500.0  |
| HR           | 75000.0  |
| Marketing    | 65000.0  |
+--------------+----------+
3 rows selected (1.244 seconds)
```

31. SELECT department, MAX(salary) FROM staff_members GROUP BY department HAVING MAX(salary) > 80000;

```
MapReduce Jobs Launched:
Stage-Stage-1:  Cumulative CPU: 0.59 sec
HDFS Read: 8992 HDFS Write: 169361115 HDFS
EC Read: 0 SUCCESS
Total MapReduce CPU Time Spent: 590 msec
+--------------+----------+
|  department  |   _c1    |
+--------------+----------+
| Engineering  | 92000.0  |
+--------------+----------+
1 row selected (1.246 seconds)
```

32. SELECT SUM(salary) FROM staff_members;

```
MapReduce Jobs Launched:
Stage-Stage-1:  Cumulative CPU: 0.24 sec
HDFS Read: 9552 HDFS Write: 169361342 HDFS
EC Read: 0 SUCCESS
Total MapReduce CPU Time Spent: 240 msec
+-----------+
|    _c0    |
+-----------+
| 317000.0  |
+-----------+
1 row selected (1.264 seconds)
```

33. SELECT MIN(salary) FROM staff_members;

```
MapReduce Jobs Launched:
Stage-Stage-1:  Cumulative CPU: 0.13 sec
HDFS Read: 10090 HDFS Write: 169361557 HDFS
 EC Read: 0 SUCCESS
Total MapReduce CPU Time Spent: 130 msec
+----------+
|   _c0    |
+----------+
| 65000.0  |
+----------+
1 row selected (1.234 seconds)
```

34. SELECT upper(full_name) FROM staff_members;

```
0: jdbc:hive2://> SELECT upper(full_name) F
ROM staff_members;
26/03/25 04:11:07 [451acfe7-4ff5-4045-ba0e-
9fe5bce20c80 main]: WARN calcite.RelOptHive
Table: No Stats for bda_assign@staff_member
s, Columns: full_name
No Stats for bda_assign@staff_members, Colu
mns: full_name
26/03/25 04:11:07 [451acfe7-4ff5-4045-ba0e-
9fe5bce20c80 main]: WARN optimizer.SimpleFe
tchOptimizer: Table bda_assign@staff_member
s is external table, falling back to filesy
stem scan.
+-----------+
|    _c0    |
+-----------+
| ABHISHEK  |
| SURENDRA  |
| AKSHAT    |
| ANJALI    |
+-----------+
4 rows selected (0.048 seconds)
```

35. SELECT length(full_name) FROM staff_members;

```
0: jdbc:hive2://> SELECT length(full_name)
FROM staff_members;
26/03/25 04:12:02 [451acfe7-4ff5-4045-ba0e-
9fe5bce20c80 main]: WARN calcite.RelOptHive
Table: No Stats for bda_assign@staff_member
s, Columns: full_name
No Stats for bda_assign@staff_members, Colu
mns: full_name
26/03/25 04:12:02 [451acfe7-4ff5-4045-ba0e-
9fe5bce20c80 main]: WARN optimizer.SimpleFe
tchOptimizer: Table bda_assign@staff_member
s is external table, falling back to filesy
stem scan.
+------+
| _c0  |
+------+
| 8    |
| 8    |
| 6    |
| 6    |
+------+
4 rows selected (0.057 seconds)
```

36. SELECT concat(full_name, ' works in ', department) FROM staff_members;

```
0: jdbc:hive2://> SELECT concat(full_name,
' works in ', department) FROM staff_member
s;
26/03/25 04:12:28 [451acfe7-4ff5-4045-ba0e-
9fe5bce20c80 main]: WARN calcite.RelOptHive
Table: No Stats for bda_assign@staff_member
s, Columns: full_name
No Stats for bda_assign@staff_members, Colu
mns: full_name
26/03/25 04:12:28 [451acfe7-4ff5-4045-ba0e-
9fe5bce20c80 main]: WARN optimizer.SimpleFe
tchOptimizer: Table bda_assign@staff_member
s is external table, falling back to filesy
stem scan.
+--------------------------------+
|              _c0               |
+--------------------------------+
| Abhishek works in Engineering  |
| Surendra works in Marketing    |
| Akshat works in Engineering    |
| Anjali works in HR             |
+--------------------------------+
4 rows selected (0.052 seconds)
```

37. SELECT current_date();

```
0: jdbc:hive2://> SELECT current_date();
26/03/25 04:12:57 [451acfe7-4ff5-4045-ba0e-
9fe5bce20c80 main]: WARN optimizer.SimpleFe
tchOptimizer: Cannot determine basic stats
for table: _dummy_database@_dummy_table fro
m metastore. Falling back.
+-------------+
|     _c0     |
+-------------+
| 2026-03-25  |
+-------------+
1 row selected (0.459 seconds)
```

38. SELECT current_timestamp();

```
0: jdbc:hive2://> SELECT current_timestamp();
26/03/25 04:13:43 [451acfe7-4ff5-4045-ba0e-9fe
5bce20c80 main]: WARN optimizer.SimpleFetchOpt
imizer: Cannot determine basic stats for table
: _dummy_database@_dummy_table from metastore.
 Falling back.
+----------------------------+
|            _c0             |
+----------------------------+
| 2026-03-25 04:13:42.834272 |
+----------------------------+
1 row selected (0.462 seconds)
```

Isha Patel
0801IT221065

39. DROP TABLE engineering_staff;

```
0: jdbc:hive2://> DROP TABLE engineering_staff;
No rows affected (0.138 seconds)
0: jdbc:hive2://> SHOW TABLES;
+----------------+
|    tab_name    |
+----------------+
| contractors    |
| staff_members  |
+----------------+
2 rows selected (0.032 seconds)
```

40. DROP DATABASE bda_assign CASCADE;

```
0: jdbc:hive2://> SHOW DATABASES;
+----------------+
| database_name  |
+----------------+
| bda_assign     |
| default        |
+----------------+
2 rows selected (0.024 seconds)
0: jdbc:hive2://> DROP DATABASE bda_assign CASCADE;
26/03/25 04:15:05 [HiveServer2-Background-Pool: Thr
ead-678]: WARN hikari.HikariConfig: txnhandler - le
akDetectionThreshold is less than 2000ms or more th
an maxLifetime, disabling it.
26/03/25 04:15:05 [HiveServer2-Background-Pool: Thr
ead-678]: WARN hikari.HikariConfig: mutex - leakDet
ectionThreshold is less than 2000ms or more than ma
xLifetime, disabling it.
26/03/25 04:15:05 [HiveServer2-Background-Pool: Thr
ead-678]: WARN hikari.HikariConfig: compactor - lea
kDetectionThreshold is less than 2000ms or more tha
n maxLifetime, disabling it.
No rows affected (0.183 seconds)
0: jdbc:hive2://> SHOW DATABASES;
+----------------+
| database_name  |
+----------------+
| default        |
+----------------+
1 row selected (0.016 seconds)
```